

Parallel Computing - Lab 4: Cuda

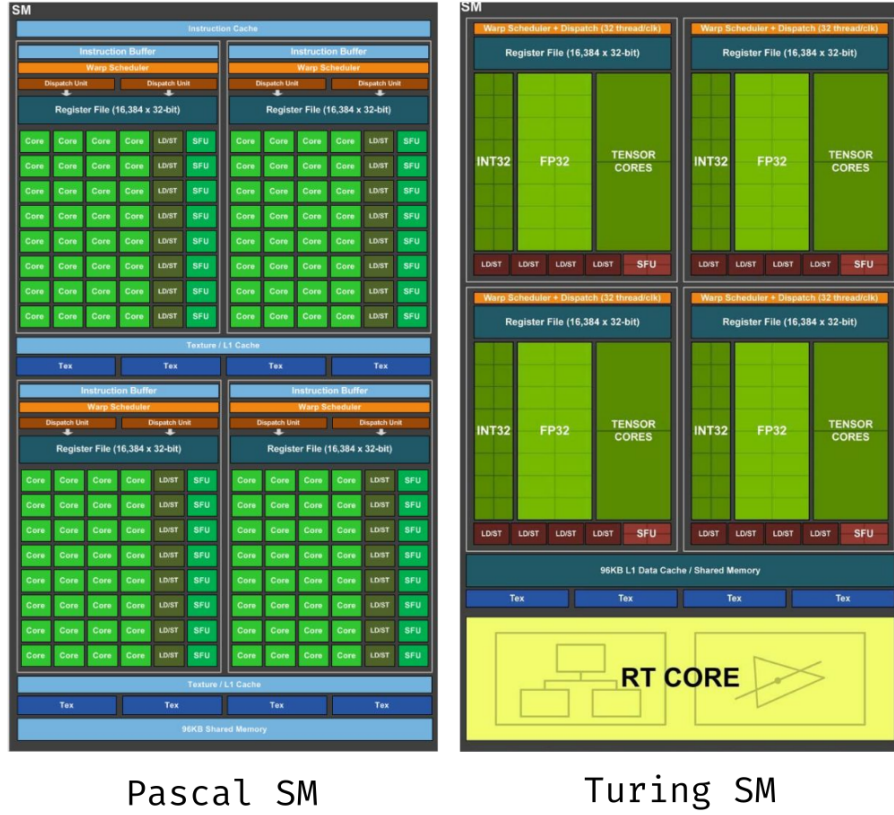
- Benjamín Ocampo: nicolasbenjaminocampo@gmail.com
- Mateo de Mayo: mateodemayo@gmail.com

Table Of Contents

1. Parallel Computing - Lab 4: Cuda
2. Hardware
3. Migration: from CPU to GPU
4. Versions
 1. Reactburst
 1. Add_source
 2. Diffuse
 3. Project
 4. Advect
 5. Kernels
 6. Pitfalls
 2. Threacburst
 3. Stepburst
 4. Migation Results
5. Failed Versions
 1. Fullburst
 2. Streamburst
 3. Graphburst
 4. Stepburst-tex
 5. Onekernel
 1. Stepburst-shmem
 2. Onekernel-shmem
 3. Stepburst-occupancy
 4. Stepburst-roae
 5. Stepburst-shidden
6. Honorable Mentions
 1. Thrust vs CUB vs ModernGPU
 2. CPU-GPU Implementation
 3. Turing Specialized Hardware
 4. CuSOLVER
 5. Loop Skewing
 6. Conclusions
 7. Speed Up Table

Hardware

- RTX 2080 Ti (target device): 68 Turing SMs
- GTX 1060 MaxQ (test device): 10 Pascal SMs



Migration: from CPU to GPU

Before putting our hands over the work and looking up optimizations, we needed to migrate all our implementation to Cuda code. It was no so easy due to the number of lines to be rewritten (more than a 1000 lines of code). Another obstacle that we had to confront was the CPU-GPU communication. Since their communication was expensive, minimizing synchronizations between them was crucial to maintain the simulation at its maximum speed. In order to do so, the entire program should be launched to the GPU at the beginning. After all the grid-updates were completed, results were returned to the CPU. We nicknamed this process as **fullburst**. That is what we were aiming. There is just one

small detail. Moving everything in a bunch leads to failures, and hours looking for their bugs. How can we move everything in that way without losing the patience?

Remember that the simulation consists of a number of updates or steps, and for each step, calls to a set of functions.

Simulation Step

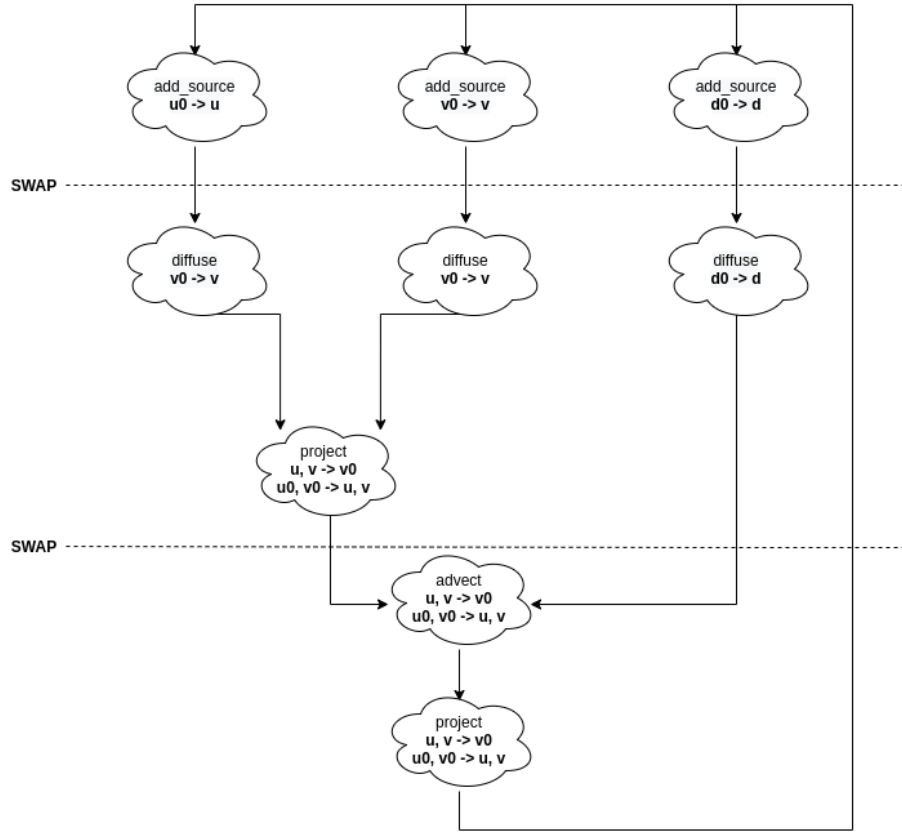
- **react**
- **step:**
 - **addsource**
 - **diffuse**
 - **project**
 - **advect**

The migration process consisted in implementing them one by one, and checking if the simulation keeps well after it. Obviously, since they were developed in an incremental way, synchronizations with the host were needed. But it was momentary until the migration was fully-implemented. It might be thought as the device “asking for help to the host”, since it can not handle the rest of the functions. We also had to take into account which snippet of code should be a **kernel**. Since one of the easiest way to synchronize threads in the GPU is through **kernels**, some dependencies were dealt with them. We are not mentioning changes of signatures, makefile rules, deletion of omp pragmas, includes, allocation of device and host data, etc. They were also developed by means of little changes but were not the core of the migration.

Versions

reactburst

During the migration, we dealt with the function **step**, i.e, implement that procedure in such a way that is executed in the device. This version was called **reactburst** since **step** was computed by the GPU, but **react** was still performed by the CPU (without changes from the lab 3). This is how the execution of **step** is:



We had to decide which cloud might be a kernel and which one might be just a host function that launches other kernels.

add_source


In this case was straightforward, since **add_source** is not a complex function and **diffuse** can not start without the previous execution of **add_source**. It was converted into a kernel.

diffuse

diffuse was split in three kernel launches. Two of them were calls to **lin_solve_rb_step** to compute Gauss-Seidel for RED and BLACK cells. These kernels had to compute just the half of the simulation matrix so the block and grid dimensions were chosen in a way that fits better for them. Since **set_bnd** depends on the previous two kernel calls, a synchronization goes between them. So another different kernel was created for its execution. Note that **set_bnd** is

launched in one dimension with different block and grid dims since it does not need to update the entire grid.

Diffuse




```

block dim = {16, 16};
grid_dim = {div_round_up(width, n) , n/block_dim.y};
sb_grid_dim = div_round_up(n + 2, block_dim.x);
sb_block_dim = block_dim.x;

for (unsigned int k = 0; k < 20; ++k) {
    gpu_lin_solve_rb_step<<<grid_dim, block_dim>>>
    gpu_lin_solve_rb_step<<<grid_dim, block_dim>>>
    gpu_set_bnd<<< sb_grid_dim , sb_block_dim>>>
}
  
```

project

project is one of the longest in terms of kernel launches, since it consists of two parts and an execution of `lin_solve`. As the previous case, they had some dependencies so project had to be split in sub kernels as above.



```

block dim = {16, 16};
grid_dim = {div_round_up(width, block_dim.x) , n/block_dim.y};
sb_grid_dim = div_round_up(n + 2, block_dim.x);
sb_block_dim = block_dim.x;

gpu_project_rb_step1<<<grid_dim, block_dim>>>(...)
gpu_project_rb_step1<<<grid_dim, block_dim>>>(...)

gpu_set_bnd<<<sb_grid_dim, sb_block_dim>>>(...)
gpu_set_bnd<<<sb_grid_dim, sb_block_dim>>>(...)

lin_solve(n, NONE, 1, 4, du0, dv0);


gpu_project_rb_step2<<<grid_dim, block_dim>>>(...)
gpu_project_rb_step2<<<grid_dim, block_dim>>>(...)

gpu_set_bnd<<<sb_grid_dim, sb_block_dim>>>(...);
gpu_set_bnd<<<sb_grid_dim, sb_block_dim>>>(...);
  
```

advect

After developing the previous two functions. `advect` was not so hard to deal with. Again the same strategy was used in this case.

advect



```
block_dim{16, 16};
grid_dim{div_round_up(width, block_dim.x), n / block_dim.y};
sb_grid_dim = div_round_up(n + 2, block_dim.x);
sb_block_dim = block_dim.x;

gpu_advect_rb<<<grid_dim, block_dim>>>
gpu_advect_rb<<<grid_dim, block_dim>>>
gpu_set_bnd<<<sb_grid_dim, sb_block_dim>>>
gpu_set_bnd<<<sb_grid_dim, sb_block_dim>>>
```

Kernels

Let us cover some of the kernels that are mentioned above. For each kernel, grid stripe loop was used. It consists in looping over the grids, working with the global id of each thread related to the current grid that we are working. Here is a part of the implementation of `gpu_lin_solve_rb_step` and `set_bnd`. The rest of kernels are similar, if you are interested in their implementation, you can glance at them at our repository.

```
__global__
void gpu_lin_solve_rb_step(grid_color color, unsigned int n, float a, float c,
                           const float *__restrict__ same0,
                           const float *__restrict__ neigh,
                           float *__restrict__ same) {
    ... // assignments
    const int grid_width = gridDim.x * blockDim.x;
    const int grid_height = gridDim.y * blockDim.y;
    const int gtidx = blockIdx.x * blockDim.x + threadIdx.x;
    const int gtidy = blockIdx.y * blockDim.y + threadIdx.y;
    for (int y = 1 + gtidy; y <= n; y += grid_height) {
        for (int x = start + gtidx; x < width - (1 - start); x += grid_width) {
            int index = y * width + x;
            same[index] = (same0[index] + a * (
                neigh[index - width] +
                neigh[index - start] +
                neigh[index - start + 1] +
                neigh[index + width]
```

```

        )) / c;
    }
}

__global__
void gpu_set_bnd(unsigned int n, boundary b, float *x) {
    const int grid_width = gridDim.x * blockDim.x;
    const int gtid = blockIdx.x * blockDim.x + threadIdx.x;
    for (unsigned int i = 1 + gtid; i <= n; i += grid_width) {
        x[IX(0, i)] = b == VERTICAL ? -x[IX(1, i)] : x[IX(1, i)];
        x[IX(n + 1, i)] = b == VERTICAL ? -x[IX(n, i)] : x[IX(n, i)];
        x[IX(i, 0)] = b == HORIZONTAL ? -x[IX(i, 1)] : x[IX(i, 1)];
        x[IX(i, n + 1)] = b == HORIZONTAL ? -x[IX(i, n)] : x[IX(i, n)];
    }
    if(gtid == 0) {
        x[IX(0, 0)] = 0.5f * (x[IX(1, 0)] + x[IX(0, 1)]);
        x[IX(0, n + 1)] = 0.5f * (x[IX(1, n + 1)] + x[IX(0, n)]);
        x[IX(n + 1, 0)] = 0.5f * (x[IX(n, 0)] + x[IX(n + 1, 1)]);
        x[IX(n + 1, n + 1)] = -0.5f * (x[IX(n, n + 1)] + x[IX(n + 1, n)]);
    }
}

```

Pitfalls

In order to complete one update of the grid, and count one step of the simulation, both `react` and `step` have to be performed. Since `react` is computed by the host we can not avoid a synchronization and copies from host to device at that point. This is what we have in `headless.c` for each simulation step.

```

react(hd_prev, hu_prev, hv_prev);
checkCudaErrors(cudaMemcpy(dd_prev, hd_prev, size_in_m, cudaMemcpyHostToDevice));
checkCudaErrors(cudaMemcpy(du_prev, hu_prev, size_in_m, cudaMemcpyHostToDevice));
checkCudaErrors(cudaMemcpy(dv_prev, hv_prev, size_in_m, cudaMemcpyHostToDevice));

step(N, diff, visc, dt,
     dd, du, dv, dd_prev, du_prev, dv_prev,
     from, to);

checkCudaErrors(cudaDeviceSynchronize());
checkCudaErrors(cudaMemcpy(hd_prev, dd_prev, size_in_m, cudaMemcpyDeviceToHost));
checkCudaErrors(cudaMemcpy(hu_prev, du_prev, size_in_m, cudaMemcpyDeviceToHost));
checkCudaErrors(cudaMemcpy(hv_prev, dv_prev, size_in_m, cudaMemcpyDeviceToHost));

```

threacburst

Obviously, the bottleneck of `reactburst` is how `react` is computed. So let us get deep into that function to see what can be done to make it faster.

```
// Two reductions
for (i = 0; i < size; i++) {
    if (max_velocity2 < uu[i] * uu[i] + vv[i] * vv[i]) {
        max_velocity2 = uu[i] * uu[i] + vv[i] * vv[i];
    }
    if (max_density < d[i]) {
        max_density = d[i];
    }
}

// Memsets
for (i = 0; i < size; i++) {
    uu[i] = vv[i] = d[i] = 0.0f;
}

// Do something with the reductions.
if (max_velocity2 < 0.0000005f) {
    uu[IX(N / 2, N / 2)] = force * 10.0f;
    vv[IX(N / 2, N / 2)] = force * 10.0f;
    for (int y = 64; y < N; y += 64)
        for (int x = 64; x < N; x += 64) {
            uu[IX(x, y)] = force * 1000.0f * (N / 2 - y) / (N / 2);
            vv[IX(x, y)] = force * 1000.0f * (N / 2 - x) / (N / 2);
        }
}

if (max_density < 1.0f) {
    d[IX(N / 2, N / 2)] = source * 10.0f;
    for (int y = 64; y < N; y += 64)
        for (int x = 64; x < N; x += 64) d[IX(x, y)] = source * 1000.0f;
}
```

Here as the comments says we have two reductions needed to compute `max_velocity2` and `max_density`. Then, the entire matrices `uu`, `vv`, and `d`, are set to 0, so `memsets` can be used here. Finally those arrays are updated so we can tell the GPU to do that for us. Instead of doing the reductions by hand, we make the most of what `thrust` (a Cuda library) and C++ have to offer.

```
// Reduction 1
dfloatp tdd_prev(dd_prev);
float max_density = *thrust::max_element(tdd_prev, tdd_prev + size);
// Implicit cudaDeviceSynchronize();

// Reduction 2
```



```

struct compare_dfloa2p2 {
    __device__
    bool operator()(tfloa2p2 lhs, tfloa2p2 rhs) {
        float lu = lhs.get<0>();
        float lv = lhs.get<1>();
        float ru = rhs.get<0>();
        float rv = rhs.get<1>();
        return lu * lu + lv * lv < ru * ru + rv * rv;
    }
};

zip_iterator<dfloa2p2> uvs_begin = make_zip_iterator(
    make_tuple(du_prev, dv_prev)
);
zip_iterator<dfloa2p2> uvs_end = make_zip_iterator(
    make_tuple(du_prev + size, dv_prev + size)
);
zip_iterator<dfloa2p2> zmaxvel2 = max_element(
    uvs_begin,
    uvs_end,
    compare_dfloa2p2()
);
// Implicit cudaDeviceSynchronize();
dfloa2p2 mv2 = zmaxvel2.get_iterator_tuple();
float mvu = *mv2.get<0>();
float mvv = *mv2.get<1>();
float max_velocity2 = mvu * mvu + mvv * mvv;

// Memsets and kernel launches
size_t size_in_mem = size * sizeof(float);
checkCudaErrors(cudaMemset(du_prev, 0, size_in_mem));
checkCudaErrors(cudaMemset(dv_prev, 0, size_in_mem));
checkCudaErrors(cudaMemset(dd_prev, 0, size_in_mem));

dim3 block_dim{16, 16};
// The gridblock mapping is one thread per reactionary point
dim3 grid_dim{
    div_round_up(div_round_up(N, 64), block_dim.x),
    div_round_up(div_round_up(N, 64), block_dim.y)
};
if (max_velocity2 < 0.0000005f)
    gpu_react_velocity<<<grid_dim, block_dim>>>(du_prev, dv_prev, force, N);

if (max_density < 1.0f)
    gpu_react_density<<<grid_dim, block_dim>>>(dd_prev, source, N);
}

```

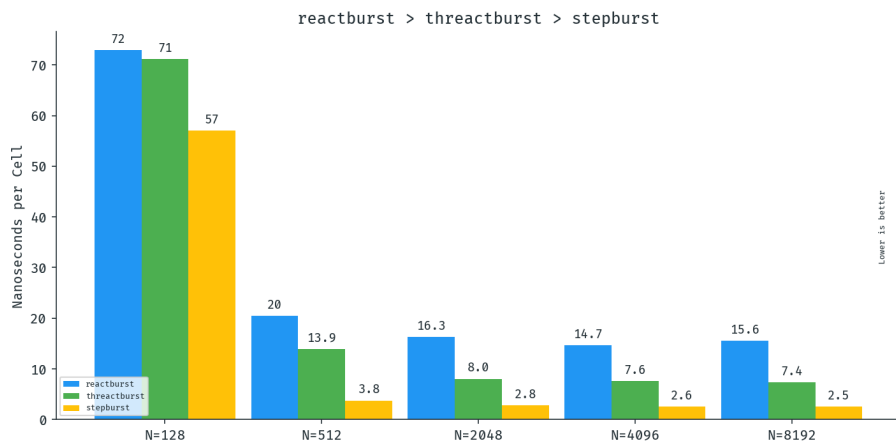
Note that in this version we are not dealing with how **react** and **step** communicate, so those copies that were shown above are still here in **threacburst**.

stepburst

It is time to make copies between **react** and **step** disappear. In this version, a bunch of kernels launches are queued up to the thrust synchronization. In order to do so, copies were moved at the end of the simulation. So everything that is touched in the program belongs to the device. This can be accomplished since the GPU works by itself, i.e, **react** and **step** use memory that is on the device. So, the CPU does not do more just launch kernels. If it were not for the thrust pitfalls this one would be the **fullburst** version that we were aiming for.

```
for (i = 0; i < steps; i++) one_step();
// Copies are placed at the end
checkCudaErrors(cudaDeviceSynchronize());
checkCudaErrors(cudaMemcpy(hd, dd, size_in_mem, cudaMemcpyDeviceToHost));
checkCudaErrors(cudaMemcpy(hu, du, size_in_mem, cudaMemcpyDeviceToHost));
checkCudaErrors(cudaMemcpy(hv, dv, size_in_mem, cudaMemcpyDeviceToHost));
checkCudaErrors(cudaMemcpy(hd_prev, dd_prev, size_in_mem, cudaMemcpyDeviceToHost));
checkCudaErrors(cudaMemcpy(hu_prev, du_prev, size_in_mem, cudaMemcpyDeviceToHost));
checkCudaErrors(cudaMemcpy(hv_prev, dv_prev, size_in_mem, cudaMemcpyDeviceToHost));
```

Migation Results



Failed Versions

The next versions can be found in their respective branches and were failed attempts at implementing ideas that in theory should've improved performance but didn't, or they were negligible, or improved in our some testing gpus (GTX 1060 MaxQ) but not in our target hardware (the RTX 2080 Ti).

fullburst

We did not want to live with that feeling of “how much performance would have been obtained if **fullburst** were implemented?” In order to check it, we did a little test that consisted in erasing the reductions in such a way thrust does not block the entire burst of launches, but still keeping **react**. In summary we just commented the following lines:

```
/*
    zip_iterator<dfloatp2> uvs_begin = make_zip_iterator(
        make_tuple(du_prev, dv_prev)
    );
    zip_iterator<dfloatp2> uvs_end = make_zip_iterator(
        make_tuple(du_prev + size, dv_prev + size)
    );
    zip_iterator<dfloatp2> zmaxvel2 = max_element(
        uvs_begin,
        uvs_end,
        compare_dfloatp2()
    );
    dfloatp2 mv2 = zmaxvel2.get_iterator_tuple();
    float mvu = *mv2.get<0>();
    float mvv = *mv2.get<1>();
    float max_velocity2 = mvu * mvu + mvv * mvv;
*/
float max_velocity2 = 0.1f;
/*
    dfloatp tdd_prev(dd_prev);
    float max_density = *thrust::max_element(tdd_prev, tdd_prev + size);
*/
float max_density = 0.1f;
```

Yep, we are cheating a bit, but the result of this trick astonished us, since we did not increase so much the performance. Why? We uncovered that the number of kernel launches that can be queued is limited to around 1024, so our idea of launching the entire program and just wait for the result is not possible. Fortunately, we discovered it before implementing it. Imagine the amount of time that we have lost if we immersed ourselves to accomplish this “optimization”.

We had already seen that cub allows reductions without blocking so that implies that react had to be implemented again.

streamburst

Another idea was the use of streams in the code and maximize concurrency. Nevertheless we did not get wildly deep into it since streams require more than a few changes in the code. What we learn after 3 laboratories is to test an idea before losing hours on a chair in front of a screen with implementations that are not going to be fruitful. The idea was to assign one stream to each kernel launch and see what happens. Something like:

```
static cudaStream_t streams[32768];
static cudaStream_t get_new_stream() {
    static int i = 0;
    cudaStreamCreate(&streams[i]);
    i++;
    return streams[i - 1];
}
```

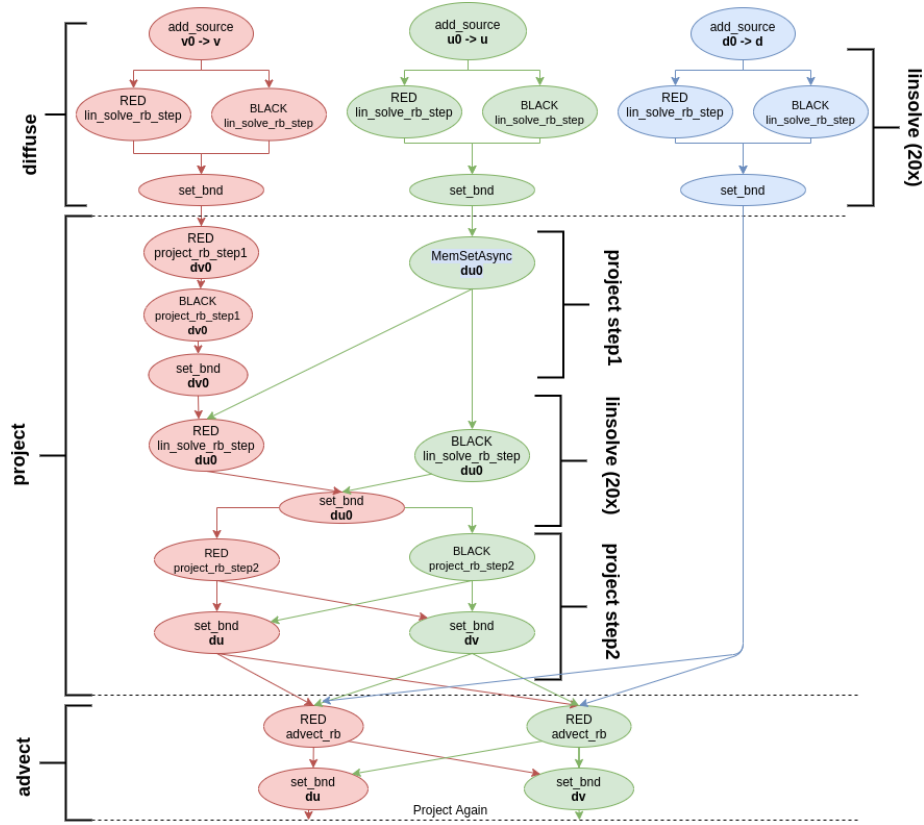
...

```
kernel<<< ..., ..., get_new_stream()>>>
```

After changing adding streams, results were oscillating between 2 and 3 ns per cell. They were not better but also not worse. So we were not sure about how much performance might be obtained with streams. Maybe there is an overhead if a new stream is created for each kernel launch.

graphburst

After doing that experiment we decide to tackle an implementation with streams but also using cuda graphs. So first we needed to analyze where the dependencies were placed in order to assign properly job/kernels for them. Looking at the graph we traced in **reactburst**, we can expand it to see kernel dependencies and find which kernel launches can be executed with other one.



As we can see in the graph, we divided the work in three streams. It might look the work is unbalanced, but consider that streams increase concurrency using resources that other kernels do not use completely. We also wanted to use cuda graphs, capturing it by means of `streamCapture` cuda directives. Since one of the conditions of these technique is to have a `main stream`, we selected the red one to command over the others.

The implementation was not so straightforward as the graph looks like, but we could do it in such a way that there are no so much changes of the `stepburst` version. At the beginning of the program, events and streams were created and then one call to step was performed enclosed by `cudaStreamBeginCapture` and `cudaStreamEndCapture`. After it, it returned an instance of graph is created and finally launched.

```
cudaGraph_t graph;
create_stream_events();
checkCudaErrors(cudaStreamBeginCapture(stream0,cudaStreamCaptureModeGlobal));
step(
    N, diff, visc, dt,
    dd, du, dv, dd_prev, du_prev, dv_prev,
```

```

    stream0, stream1, stream2,
    spread, join_stream0, join_stream1, join_stream2
);
checkCudaErrors(cudaStreamEndCapture(stream0, &graph));
checkCudaErrors(cudaGraphInstantiate(&step_graph, graph, NULL, NULL, 0));
cudaGraphDestroy(graph);

one_step(...){
    react(...);
    checkCudaErrors(cudaGraphLaunch(step_graph, stream0));
    ...
}

```

Unfortunately it seems that using streams and cuda graphs was not the key of the problem since we could not get improvements. In order to see if streams were actually used we check it with nvvp that show us a thoroughly execution of the program. Here we could obtain that `gpu_lin_solve_rb_step` is the 85% of the code, and the rest is taken by the other kernels. Maybe we could get some concurrency between kernels but they were not the key of the problem.

stepburst-tex

The cuda programming guide and lots of posts praise the use of texture memory instead of reading from global memory. Especially for stencil patterns. Nevertheless its implementation would be a little painful for us, since textures must be declared in file scope (In order to be known at compile time), we can not pass them by parameter. That would lead to so many changes in functions signatures and hours of testing correctness. Instead of that, we test it with the kernel `gpu_lin_solve_rb_step` checking if they speed up that kernel.

```

__global__
void gpu_lin_solve_rb_step_tex(...){
    ...
    for (int y = 1 + gtidy; y <= n; y += grid_height) {
        for (int x = start + gtidx; x < width - (1 - start); x += grid_width) {
            int index = y * width + x;
            same[index] = (same0[index] + a * (
                tex2D(tex_dred, x, y - 1) +
                tex2D(tex_dred, x - start, y) +
                tex2D(tex_dred, x - start + 1, y) +
                tex2D(tex_dred, x, y + 1)
            )) / c;
        }
    }
}

```

```

texture<float, cudaTextureType2D, cudaReadModeElementType> tex_dred;

// Bind the texture @tex_dred to the array @dred.
cudaChannelFormatDesc channel_desc = cudaCreateChannelDesc<float>();
checkCudaErrors(cudaBindTexture2D(NULL, &tex_dred, dred, &channel_desc, width, n, width * s);

// Gauss-Seidel
for (unsigned int k = 0; k < 20; ++k) {
    // Red case is launched both since we can not pass textures by parameter
    gpu_lin_solve_rb_step_tex<<<grid_dim, block_dim>>>(RED, ...)
    gpu_lin_solve_rb_step_tex<<<grid_dim, block_dim>>>(RED, ...);
    gpu_set_bnd<<<div_round_up(n + 2, block_dim.x), block_dim.x>>>(n, b, dd);
}
cudaUnbindTexture(tex_dred);

```

Then using `nvprof` we could measure the time needed to run this kernel with the older one. Unfortunately for us, results were so similar. We discovered later that Turing merges the cache L1 and the texture cache in such a way that when we launch kernels both caches are used. We also wanted to use textures in `advect` to take the most of `interpolations`. But we could not use it due to red-black organization. The idea was to bind a matrix `dd`, `du`, and `dv`, to a texture and avoid parsing float coordinates to indexes. Nevertheless, we need the indexes to know if a cell is RED or BLACK.

onekernel

One of our main concerns was the cache not being properly utilized as each `lin_solve` kernel launch erases any cache the previous blocks may have fetched, and by the nature of the problem, this would imply to re-read twenty times from global memory.

To prevent this, in recent versions of CUDA the feature of Cooperative Groups (CG) was introduced. This are groups of threads that can be smaller than a warp, bigger than an SM covering an entire device worth of threads and even bigger than a device extending over multiple devices. CG offer lots of possibilities that are discussed more in depth in **this article** but for our particular case, what we needed was a way to synchronize the entire device between `gpu_lin_solve_rb_step` launches, and the cooperative groups `.sync()` method does exactly that. So in this way we could make `gpu_lin_solve` to be a single kernel that iteratively and in a synchronized manner, calls `gpu_lin_solve_rb_step` device functions and syncs them with `grid.sync()` as follows:

```

__device__ gpu_lin_solve_rb_step(...);
__global__ gpu_lin_solve(...) {
    for (int k = 0; k < 20; k++) {
        gpu_lin_solve_rb_step(RED, ...);
    }
}

```

```

        gpu_lin_solve_rb_step(BLACK, ...);
        grid.sync();
    }
}

```

Unfortunately it seems that the **grid synchronization feature** of the CG is not very optimized in hardware yet (at least up to turing), and unfortunately we saw a decrease in performance (from ~2.5 to ~3.5 with N=4096).

Some points that we need to mention that appeared during this implementation were:

- Not syncing: Removing the `grid.sync()` call above makes the simulation a little weird, however it could be still a pretty reasonable approximation for some use cases, and the performance increases significantly (from ~2.5 to ~1.6 with N=4096).
- Occupancy API: For using the grid sync feature of CGs we need to use special ways of launching kernels and defining their dimensions. In particular the number of blocks must match those of the amount of SMs in the device, and the number of threads in a block must be calculated with the some of the Occupancy API methods.
- Cache Preference: In this example we ought to use the `cudaDeviceSetCacheConfig` function which allows us to hint to the device that we prefer Cache L1 over shared memory as we are not using none of the latter. However in all of our tests in this and other implementations it seems that the device is smart enough to tell which setting to use as the performance doesn't improve.

stepburst-shmem

A lot of places talked wonders about the usage of shared memory even in trivial **examples** where the cache L1 should work perfectly fine so we decided to give it a try.

```

__device__
void gpu_lin_solve_rb_step(...) {
    __shared__ float neigh_cache[BLOCK_WIDTH + 2][BLOCK_HEIGHT + 1];
    ... // load neighbours to cache
    __syncthreads();
    same[...] = ... // something related to same0 and neigh_cache but not neigh
}

```

Later on we profiled with `nvcc -m global_hit_rate` and discovered a hit rate of 40%. With this and by commenting and discommenting reads, we tracked down the following pattern of cache usage (or at least our hypothesis for it):

Access	Referenced from
Access	Referenced from
same	write miss?
same0	miss
neigh up	hit
neigh down	hit
neigh left	miss
neigh right	vectorized

So this explains our 40% (2/5) accesses as one of them is vectorized. Then we did an experiment and not a full implementation of the same idea but putting same0 in shared as it would've been easier for not having to deal with the neighbours halo, but still it did not give gains.

onekernel-shmem

As in this version we have exactly one block per SM per the CG requirements, the idea is to fill the shared memory of each block with a portion of its sub grid and use shared memory when the grid strid loop needs to use memory that is cached there. Unfortunately in our tests the with and without the added conditional were in all cases worst than stepburst. The grid synchronization was still a heavy hit on performance.

stepburst-occupancy

We brought the idea of using the occupancy api from the onekernel implementation with the hopes that by dividing the problem to match exactly the hardware it may improve the performance, but the opposite happened. It seems that dividing the problem in weird partitions (in this cases divided by 68 as the RTX 2080 ti has 68 SMs) is probably sub utilizing the vector units.

stepburst-roae

As we were running out of good ideas, we implemented all of our bad ones in here, the “Root Of All Evil” branch, which applies and discuss lots of minimal and mostly superficial optimizations, which are as follows:

- `cudaDeviceSetCacheConfig`: Hint for increased cache usage as we couldn't use shared memory

- Vector loads: `float4`, `int4`, and similar types are supposed to increase a bit **the overall bandwidth** but for using them we needed to make aligned reads and our memory layout did not allow that.
- `#pragma loop unroll`: For unrolling loops.
- `nvcc` flags: basically `-use_fast_math` and `friends`.
- Block/grid dims: There were many settings which gave good results, we decided to go with 8x8 as 64 is the amount of cuda cores in one SM for the 2080 Ti.
- PTX intrinsics: There are many intrinsics for math (which were already enabled by the `nvcc` flags) and for memory load/store, we tried using a similar technique as to the `stream` technique we used in lab2 with `__stwt` but it didn't give any gains, probably CUDA was already figuring that out. Also the `__ldg` intrinsic which hints that a load should be cached in a read-only faster cache, it improved about 2% the performance in our pascal GTX 1060 MaxQ but nothing on the target RTX 2080 Ti.

stepburst-shidden

So this was the hacky idea. We read in some places that the shared memory even though in theory it has the same lifetime of a block, in practice the driver does not clean shared memory between calls and just sets it unused. So the idea is that we can reuse between kernel launches this “undefined” shared memory with some tricks which we implemented as follows:

```
void lin_solve(...) {
    gpu_lin_solve_rb_step_shtore(RED);
    gpu_lin_solve_rb_step_shtore(BLACK);
    for (int k = 0; k < 19; k++) {
        gpu_lin_solve_rb_step_shload(RED);
        gpu_lin_solve_rb_step_shload(BLACK);
    }
}

__global__
void gpu_lin_solve_rb_step_shtore(...) {
    __shared__ float csame0[BLOCK_HEIGHT][BLOCK_WIDTH];
    __shared__ int bx, by, id1, id2;

    // Save block information for later discovery
    if (first thread of block) {
        bx = blockIdx.x; by = blockIdx.y;
        id1 = bx * by + bx; id2 = bx * by + by;
    }

    // Save appropriate value of same0 for next kernels
```

```

const float previous = same0[index];
csame0[threadIdx.y][threadIdx.x] = previous;

... // Get lots of "set but not used" compiler warnings

... // regular lin_solve_rb_step using previous instead of same0
}

__global__
void gpu_lin_solve_rb_step_shload(...) {
    __shared__ float csame0[BLOCK_HEIGHT][BLOCK_WIDTH];
    __shared__ int bx, by, id1, id2;

    float previous;
    // Notice most of these variables are undefined
    if (id1 == bx * by + bx && id2 == bx * by + by) {
        // Found'em! read "undefined" memory.
        previous = csame0[threadIdx.y][threadIdx.x];
    } else {
        // Read from global memory
        previous = same0[index]; // (*) There is a catch here, explained in a moment
    }

    ... // Regular lin_solve_rb_step but using bx and by as block indexes instead
}

```

Notice in the conditional from `_shload` we check that our undefined shared memory effectively has the correct values, one could ask about the ratio in which this check effectively “finds” what it is looking for, here are the results:

GPU	$N \leq 64$	$N = 512$	$N = 2048$
GTX 1060 MaxQ	100%	50%	50%
RTX 2080 Ti	100%	25%	2.21%

So from one side, it seems that the new turing architecture does something with undefined shared memory that makes it not as reliable as in pascal. In pascal however that 50% means literally half of the reads to `same0` in the twenty iterations could be avoided.

(*) Another thing to take into consideration is the `else` in the code above is figurative and does not solve all of our problems to make the algorithm correct. In particular, when we don’t find a hit, we can’t just keep computing as usual, because our sub grid could’ve been already computed by another block which found the corresponding shared memory. And as such we would need to divide

into two steps the solver, first filling all the found-in-undefined-shared-memory sub grids, and then all the rest which were not cached.

So because of that implementation complication and the fact that in turing, our target architecture, would not make a notorious difference, we decided not to finish implementing the idea.

Honorable Mentions

During the development with CUDA, lots of appealing options and names appeared and so we thought it could be a good idea to mention some of them.

Thrust vs CUB vs ModernGPU

These are three libraries that you may encounter while looking for cuda helper libraries, the focus of each one is different and so you should know which one to look for.

- Thrust: focuses on ease of use, and high level abstractions that run not only on CUDA but also on OpenMP, OpenACC and others. Thrust uses CUB at its core.
- CUB: on the other hand is a specific, performance oriented library for CUDA, it is more cumbersome to use but it is more flexible as well. We particularly notice the ability of CUB to do reductions on device while thrust forces you to bring the result to host.
- ModernGPU: Is more of a learning oriented library that is made to be read and not so much used, the codebase prioritize ease of understanding and novel ideas to get inspired from. It is a go to if you are trying to solve a problem in a novel way.

CPU-GPU Implementation

We did not get to implement an heterogenous version which uses both OpenMP and CUDA, but we thought about it a lot. The rough idea of implementation would've been to reuse our OpenMP strip-divided version delimiting the grid with a particular row index. In each step the CPU and GPU need to interchange the topmost (or bottommost) row to the other and continue its work.

Turing Specialized Hardware

Turing implements **two pieces** of novel hardware that we considered to use but did not really found a good application for them to dig into an test implementation.

- Tensor Cores: Are special functional units in each SM (8 of them in each SM) dedicated to 4x4 matrix fused multiply and add.
- RT Cores: Are a piece of dedicated hardware found in each SM that solves two problems related to ray tracing: ray casting (intersection test) and boundary volume hierarchy traversal (BVH).


cuSOLVER

cuSOLVER is a library that implements many linear solvers in CUDA, unfortunately there was no explicit “Gauss-Seidel Relaxation” solver and our numerical analysis knowledge did not help much helping us to recognize which solver would’ve been of use in our case.

Loop Skewing

Loop or Time skewing algorithms for solving stencil problems are what **this paper** calls the *Wavefront Pattern* and it was a concept we bumped into many times on our searches. We never got to implement it but the idea is relatively simple:


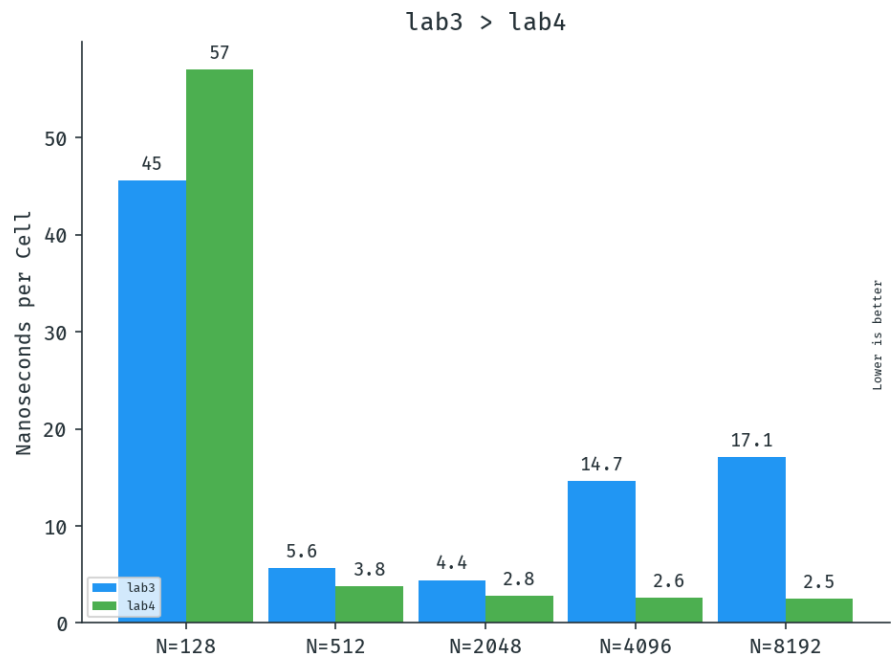
In our gauss-seidel traversal, if you consider the dependencies of each cell, you will notice that diagonal cells do not have any dependency between them, and as such they can be computed in parallel, having each iteration of a parallel loop traverse (*like wavefronts*) from the upper left cell in diagonals to the bottom right cell in the following way is a correct parallel implementation.




Conclusions

Here we can see all our results during these course. Before CP we would not have considered the use of vectorization, omp, cuda, and even compiler flags. So we thought the best way to close the subject would be compare all our bests tricks and optimizations. Note that in the comparison between lab3 and lab4,

we got better results with cheaper hardware. This is due to the high bandwidth that the GPU offers us.





Speed Up Table

Speed up with respect from the previous version (starting with the -00 implementation)

N	lab1	lab2	lab3	lab4	total
128	1.05x	10.8x	2.04x	0.79x	18x
512	1.46x	11.1x	15.7x	1.47x	375x
2048	1.69x	11.4x	20.0x	1.57x	608x
4096	2.32x	10.2x	6.87x	5.65x	916x
8192	2.27x	9.85x	6.14x	6.84x	940x